

Flower Type Recognizer

Imports

```
In [1]: %matplotlib inline
%config InlineBackend.figure_format = 'retina'

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow_datasets as tfds
import tf_keras
import logging
import json
import tensorflow_hub as hub
import pandas as pd
```

WARNING:tensorflow:From C:\Users\HP\PycharmProjects\pythonProject\venv\Lib\site-packages\tf_keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [2]: tfds.disable_progress_bar()
```

```
In [3]: logger = tf.get_logger()
logger.setLevel(logging.ERROR)
physical_devices = tf.config.list_physical_devices('GPU')
for device in physical_devices:
    tf.config.experimental.set_memory_growth(device, True)
```

Load the Data

- Use Tensorflow dataset `oxford_flowers102`
- split it into
 - train
 - test
 - validation

```
In [4]: # Split the data to train, val, test as the splits the original data provides
splits = ['train', 'validation', 'test']

# Load the 'oxford_flowers102' dataset and split it into `splits`
dataset, dataset_info = tfds.load('oxford_flowers102', as_supervised=True, with_info=True)

validation_set, test_set, training_set = dataset
len(training_set), len(validation_set), len(test_set)
```

Out[4]: (6149, 1020, 1020)

Explor the Dataset

```
In [5]: dataset_info.features['image']
```

Out[5]: Image(shape=(None, None, 3), dtype=uint8)

```
In [6]: dataset_info.features['label']
```

Out[6]: ClassLabel(shape=(), dtype=int64, num_classes=102)

```
In [7]: total_examples = dataset_info.splits['train'].num_examples + dataset_info.splits['t
num_training_examples = len(training_set)
num_validation_examples = len(validation_set)
num_test_examples = len(test_set)

print(f'There are {total_examples:,} images in Total')
print(f'There are {num_training_examples:,} images in the training set')
print(f'There are {num_validation_examples:,} images in the validation set')
print(f'There are {num_test_examples:,} images in the test set')
```

There are 8,189 images in Total

There are 6,149 images in the training set

There are 1,020 images in the validation set

There are 1,020 images in the test set

Data Info Summary

Image:

- **Shape:** (None, None, 3)
 - **dtype:** uint8
-

Label:

- **Number of Classes:** 102
 - **dtype:** int64
-

Examples

- There are **7,169** images in Total
- There are **816** images in the training set
- There are **204** images in the validation set

- There are **6,149** images in the test set

Exploring The Images

Print first three images shape

```
In [8]: i = 1 # For counting which image we are in currently
        for image, label in training_set.take(3):
            print(f'Image_{i} Shape: {image.shape}, Label: {label}')
            i += 1
```

Image_1 Shape: (542, 500, 3), Label: 40

Image_2 Shape: (748, 500, 3), Label: 76


Image_3 Shape: (500, 600, 3), Label: 42

Print first image

```
In [9]: for image, label in training_set.take(1):
        image = image.numpy().squeeze() # to get rid of `1`
        label = label.numpy()

        plt.figure(figsize=(2,1))
        plt.imshow(image)
        print(f"Label is {label}")
```

Label is 40



Attach the labels to theirs corresponding names

```
In [10]: with open('label_map.json', 'r') as f:
        class_names = json.load(f)
        for key, value in class_names.items():
            print(f"Label: `{key}` -- Name: `{value}`")
            if key == '5':
                break
```

Label: `0` -- Name: `pink primrose`

Label: `1` -- Name: `hard-leaved pocket orchid`

Label: `2` -- Name: `canterbury bells`


Label: `3` -- Name: `sweet pea`

Label: `4` -- Name: `english marigold`

Label: `5` -- Name: `tiger lily`

```
In [11]: for image, label in training_set.take(1):
          image = image.numpy().squeeze()
          label = str(label.numpy())
          plt.figure(figsize=(2,1))
          plt.imshow(image)
          print(f"This is an image of {class_names[label]}")
```

This is an image of barbeton daisy



prepare The Dataset

Steps:

- Shuffle the data (shuffle each quarter for memory can handle)
- normalize the data
- make batches (size 64)
- prefetch (prepare the next batch while the first is execution)

```
In [12]: batch_size = 64
          image_size = 224

          def normalize(image, label):
              image = tf.cast(image, tf.float32)
              image = tf.image.resize(image, (image_size, image_size)) #resize to (224,224,3)
              image /= 255 # (normalize the range of 0-255 -> 0-1)
              return image, label

          training_batches = training_set.shuffle(num_training_examples//4).map(normalize).ba
          validation_batches = validation_set.shuffle(num_validation_examples//4).map(normali
          testing_batches = test_set.shuffle(num_test_examples // 4).map(normalize).batch(bat
```

Use A Pre-trained Module

Steps:

- Load the module
- Freeze the weights
- Decide the input shape

```
In [13]: mobilenet_URL = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4

# Load the MobileNet model from TensorFlow Hub
mobilenet_model = hub.KerasLayer(mobilenet_URL, input_shape=(224, 224, 3), trainable=True)

In [14]: # Freeze the weights, small data set with similar data.
mobilenet_model.trainable = False

In [15]: from tf_keras.utils import plot_model
```

Build the network

Steps:

- Connect the old network
- **Layer1:** Activation: `relu` - Height: `256`
- **Layer2:** Activation: `relu` - Height: `128`
- **Output Layer:** Activation: `Softmax` - Height: `102`

```
In [18]: import tf_keras

number_of_classes = len(class_names)

model = tf_keras.Sequential([
    mobilenet_model,
    tf_keras.layers.Dense(256, activation='relu'),
    tf_keras.layers.Dense(128, activation='relu'),
    tf_keras.layers.Dense(number_of_classes, activation='softmax')
])

model.summary()
```

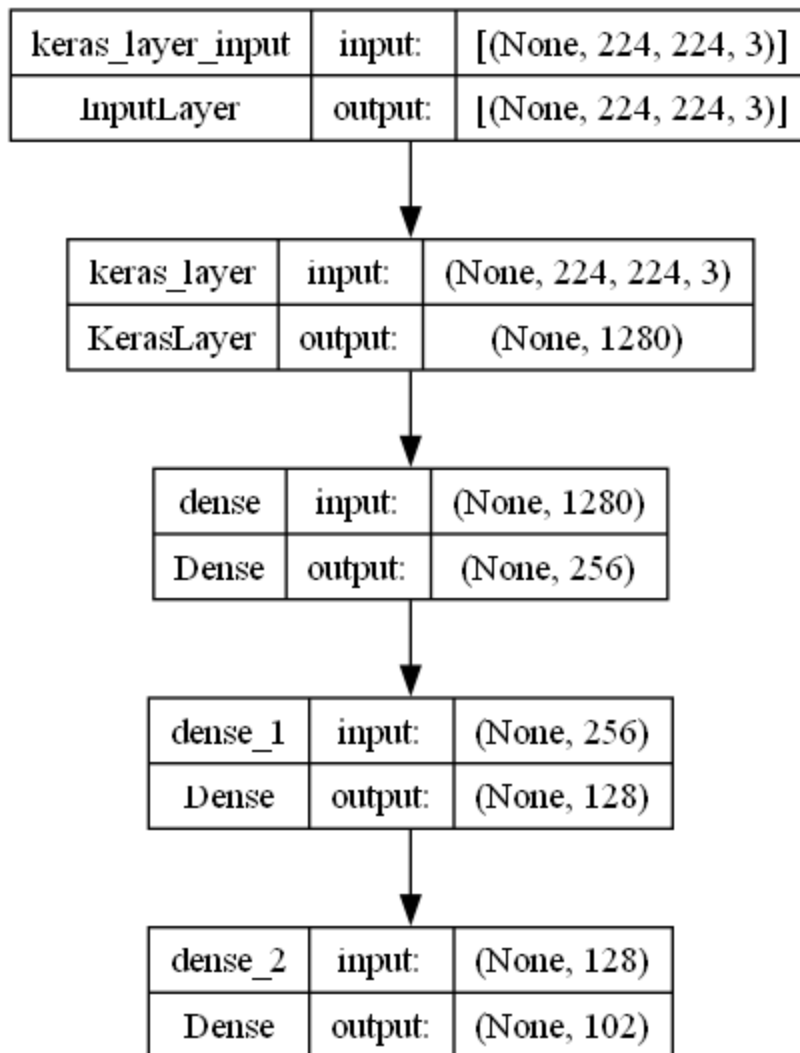
Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|--------------|---------|
| ===== | | |
| keras_layer (KerasLayer) | (None, 1280) | 2257984 |
| dense_3 (Dense) | (None, 256) | 327936 |
| dense_4 (Dense) | (None, 128) | 32896 |
| dense_5 (Dense) | (None, 102) | 13158 |
| ===== | | |
| Total params: 2631974 (10.04 MB) | | |
| Trainable params: 373990 (1.43 MB) | | |
| Non-trainable params: 2257984 (8.61 MB) | | |

Network plot

```
In [17]: plot_model(model, to_file="mobilenet_model.png", show_shapes=True, show_layer_names
```

Out[17]:



Training The Module

Fetures:

- optimizer = adam loss = SGD
- Patience: 5 (validation accuracy loss)

```
In [19]: model.compile(optimizer='adam',  
                        loss='sparse_categorical_crossentropy',  
                        metrics=['accuracy'])
```



```
In [20]: EPOCHS = 12  
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)  
  
history = model.fit(training_batches,
```

```
epochs=EPOCHS,  
validation_data=validation_batches,  
callbacks=[early_stopping])
```

```
Epoch 1/12  
97/97 [=====] - 75s 680ms/step - loss: 2.4244 - accuracy:  
0.4661 - val_loss: 1.5335 - val_accuracy: 0.6176  
Epoch 2/12  
97/97 [=====] - 57s 576ms/step - loss: 0.7066 - accuracy:  
0.8193 - val_loss: 0.8677 - val_accuracy: 0.7716  
Epoch 3/12  
97/97 [=====] - 65s 659ms/step - loss: 0.3660 - accuracy:  
0.9102 - val_loss: 0.7033 - val_accuracy: 0.8000  
Epoch 4/12  
97/97 [=====] - 56s 573ms/step - loss: 0.2235 - accuracy:  
0.9470 - val_loss: 0.6148 - val_accuracy: 0.8196  
Epoch 5/12  
97/97 [=====] - 55s 564ms/step - loss: 0.1328 - accuracy:  
0.9725 - val_loss: 0.5898 - val_accuracy: 0.8353  
Epoch 6/12  
97/97 [=====] - 64s 651ms/step - loss: 0.0879 - accuracy:  
0.9823 - val_loss: 0.5531 - val_accuracy: 0.8412  
Epoch 7/12  
97/97 [=====] - 62s 632ms/step - loss: 0.0556 - accuracy:  
0.9919 - val_loss: 0.5044 - val_accuracy: 0.8471  
Epoch 8/12  
97/97 [=====] - 61s 617ms/step - loss: 0.0370 - accuracy:  
0.9959 - val_loss: 0.5149 - val_accuracy: 0.8520  
Epoch 9/12  
97/97 [=====] - 58s 592ms/step - loss: 0.0263 - accuracy:  
0.9982 - val_loss: 0.4917 - val_accuracy: 0.8559  
Epoch 10/12  
97/97 [=====] - 68s 690ms/step - loss: 0.0168 - accuracy:  
0.9995 - val_loss: 0.4845 - val_accuracy: 0.8676  
Epoch 11/12  
97/97 [=====] - 68s 695ms/step - loss: 0.0104 - accuracy:  
0.9998 - val_loss: 0.4966 - val_accuracy: 0.8608  
Epoch 12/12  
97/97 [=====] - 63s 645ms/step - loss: 0.0097 - accuracy:  
0.9998 - val_loss: 0.5043 - val_accuracy: 0.8647
```

```
In [21]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
plt.figure(figsize=(10, 4))  
plt.plot(acc, label='Training Accuracy')  
plt.plot(val_acc, label='Validation Accuracy')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.title('Training & Validation Accuracy')  
plt.legend()  
plt.show()
```

```
plt.figure(figsize=(10, 4))
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training & Validation Loss')
plt.legend()
plt.show()
```



Testing The module

```
In [22]: loss, accuracy = model.evaluate(testing_batches)

print(f"Loss: {loss * 100}%\nAccuracy: {accuracy * 100}%")
```


16/16 [=====] - 10s 572ms/step - loss: 0.4033 - accuracy: 0.8853
Loss: 40.3276264667511%
Accuracy: 88.52941393852234%

Save The Module

```
In [23]: saved_keras_model_filepath = f'./Flower_Recognizer.h5'

model.save(saved_keras_model_filepath)
```

C:\Users\HP\PycharmProjects\pythonProject\venv\Lib\site-packages\tf_keras\src\engine\training.py:3098: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native TF-Keras format, e.g. `model.save('my_model.keras')`.
saving_api.save_model(

```
In [24]: !unzip -t {saved_keras_model_filepath}
```

'unzip' is not recognized as an internal or external command,
operable program or batch file.

```
In [25]: reloaded_keras_model = tf_keras.models.load_model(
        saved_keras_model_filepath,
        custom_objects={'KerasLayer': hub.KerasLayer}
    )
```

```
In [33]: (model.get_weights()[0][0][0] == reloaded_keras_model.get_weights()[0][0][0][:5]
```

```
Out[33]: array([[ True],
               [ True],
               [ True],
               [ True],
               [ True]])
```

Some tests on the module

1- Prepare The Image:

- resize to what module take (224,224,3)
- normalize from (0-255) -> (0-1)

```
In [27]: pre_image_size = 224

def process_image(image):
    image = tf.convert_to_tensor(image, dtype=tf.float32)
    image = tf.image.resize(image, (pre_image_size, pre_image_size))
    image /= 255
    image = image.numpy()
    return image
```


2- Print The Processed Image

```
In [28]: from PIL import Image

image_path = './test_images/hard-leaved_pocket_orchid.jpg'
im = Image.open(image_path)
test_image = np.asarray(im)

processed_test_image = process_image(test_image)

fig, (ax1, ax2) = plt.subplots(figsize=(10,10), ncols=2)
ax1.imshow(test_image)
ax1.set_title('Original Image')
ax2.imshow(processed_test_image)
ax2.set_title('Processed Image')
plt.tight_layout()
plt.show()
```



3- Make Predict Function

```
In [29]: def predict(image_path, model, k):
    image = Image.open(image_path)
    image = np.asarray(image)
    image = process_image(image)
    image = np.expand_dims(image, axis=0)

    prediction_dict = model.predict(image)
    pandas_predictions = pd.DataFrame(prediction_dict)
    sorted_prediction = pandas_predictions.T.sort_values(by=0, ascending=False).head(k)
    probs = sorted_prediction.values[0]
    classes = sorted_prediction.keys().tolist()
    # print(sorted_prediction)
    return list(probs), list(classes)
```

4- Function to plot the images

```
In [30]: def plot_image_probs(image_path, model= model, k=5):
    probs, classes = predict(image_path, model, k)
    names = [class_names[f"{i}"] for i in classes]

    fig, (ax1, ax2) = plt.subplots(figsize=(6, 9), ncols=2)

    ax1.imshow(Image.open(image_path))
    ax1.set_title('Original Image')
    ax1.axis('off')

    ax2.barh(names, probs, color='blue')
    ax2.set_aspect(0.1)
    ax2.set_yticks(np.arange(len(names)))
    ax2.set_yticklabels(names)
    ax2.set_title('Class Probability')
    ax2.set_xlim(0, 1.1)

    plt.tight_layout()
    plt.show()
```


5- Test The photos on test_images

```
In [31]: import os

    directory = "./test_images"
    for i in os.listdir(directory):
        path = directory + "/" + i
        plot_image_probs(path, model, 5)
```


1/1 [=====] - 2s 2s/step

Original Image





1/1 [=====] - 0s 53ms/step

Original Image



1/1 [=====] - 0s 50ms/step

Original Image



1/1 [=====] - 0s 50ms/step

Original Image

